

```
1 Homework #2 for protok11
2 Peter Boström <pbos@kth.se>
3 2011-09-11
4
5 1. ARP
6
7 Hosts:
8
9 H1: IP=A, MAC=x
10 H2: IP=B, MAC=y
11 H3: IP=C, MAC=z
12 H4: IP=D, MAC=t
13
14 Bridge B1 Connects H3, H4 and router R1.
15
16 Router R1:
17
18 => B1: IP=E, MAC=v
19 => H2: IP=F, MAC=w
20 => H1: IP=G, MAC=u
21
22 B1 has a MAC address table, R1 and all hosts have ARP tables, all initially
23 empty.
24
25 Consider that host H4 sends an IPv4 unicast datagram to host H1.
26
27     a, b) Provide the state of the five ARP caches as they will appear after the
28           IPv4 unicast datagram has been delivered to host H1, that is, after
29           dynamic ARP resolution has been made.
30
31     H4 checks its routing table and decides to send this to IP E (R1's IP on
32     the interface connected to the bridge). E is not in H4's ARP table, so it
33     broadcasts a request for E's MAC address through ARP. This request includes
34     H4's IP and MAC address. B1 learns H4's MAC and repeats the broadcast on its
35     other interfaces. H3 learns H4's IP and MAC address, but sends nothing. R1
36     learns H4's IP and MAC as well, and send back its own IP and MAC address
37     directly to H4. B1 learns R1's MAC and interface. H4 learns E's MAC address.
38
39     H4 is now ready and forwards the datagram to R1. R1 checks its routing table
40     and decides to send the datagram to H1's IP A. It broadcast to find A's MAC
41     address to be able to deliver it. By doing so it also sends its own IP and
42     MAC, which gets learned by H1 before returning (A, x) to R1.
43
44     The datagram then gets delivered to H1.
45
46     a) ARP caches
47
48     H1 (IP, MAC):
49     G, u
50
51     H3 (IP, MAC):
52     D, t
53
54     H4 (IP, MAC):
55     E, v
56
57     R1 (IP, MAC):
58     D, t
59     A, x
60
61     H2 is empty.
62
63     b) MAC address table
64
65     B1 (MAC, Interface):
66     t, West
67     v, South
```

```
68
69 2. UDP and fragmentation
70
71 6400 octets to be sent inside an UDP datagram over a link with MTU=2272 octets.
72
73 a) How many fragments are transmitted?
74
75 An UDP datagram adds an 8-octet header => 6408 bytes. An IPv4 header will
76 take 20 bytes which gives a maximum payload of 2252. Fragmented packets must
77 however be in sizes divisible by 8, so we can use up to 2248 octets for each
78 packet.
79
80 We'll need  $\text{ceil}(6408/2248) = 3$  fragments to send this UDP datagram.
81
82 b) Give the values of the MF bit, offset and total length field of the IP
83 header of each fragment.
84
85 Packet 1 will have MF=1 and offset = 0, it'll contain the first 2248 octets
86 of the UDP datagram. Length = 2248. There are now 4160 octets left.
87
88 Packet 2 will have MF=1 and offset =  $2248/8 = 281$  and contain the 2248
89 following octets. Length = 2248. There are now 1912 octets left.
90
91 Packet 3 will have MF=0 and offset =  $2248*2/8 = 562$  and contain the remaining
92 1912 octets. Length = 1912. This is the final packet.
93
94 3. TCP session management
95
96 Consider the TCP SYN Flooding Attack.
97
98 a) What kind of attack is this?
99
100 A Denial-of-Service attack.
101
102 b) How is the attack done?
103
104 During a TCP handshake, a SYN is sent to the server, which sends a SYN-ACK
105 back the attacker is expected to send an ACK back, which never happens. This
106 gives the server a half-open connection and occupies limited resources.
107
108 An attacker will send a lot of these SYN messages and establish a lot of
109 half-open connections, eventually the SYN queue or other limited resources
110 will be occupied and connections can't be established by legitimate hosts.
111
112 c) What damage is caused?
113
114 New hosts will be unable to connect to the server, as it's occupied with these
115 half-open connections.
116
117 d) How can a server alleviate the effects of the attack? Describe one
118 solution in sufficient detail such that it becomes clear why it alleviates
119 the attack.
120
121 Provided that the attacker is unable to spoof the source address, then you can
122 simply limit the number of new connections per source per timeframe. This means
123 that additional SYN messages from the same source are dropped, and the number
124 of half-open connections from that host is limited.
125
126 Another technique is known as SYN cookies which presents a cryptographic
127 task that the source needs to solve.
128
129 4. TCP 2
130
131 Segment to existing connection is sent at 4:30:20. The sender doesn't receive an
132 acknowledgement and resends at 4:30:28 and receives an acknowledgement at 4:30:30
133 where it sends another segment and receives the following 2 seconds after, at 4:30:32.
134 When the first mentioned segment was sent the smoothed RTT (sRTT) was 4 seconds.
```

```

135
136 Give the values of the smoothed RTT (sRTT), the variation (RTTvar) and the
137 retransmission timeout (RT0):
138
139 RFC2988:
140
141 (2.3) When a subsequent RTT measurement R' is made, a host MUST set
142
143     RTTvar <- (1 - beta) * RTTvar + beta * |sRTT - R'|
144     sRTT <- (1 - alpha) * sRTT + alpha * R'
145
146 The value of sRTT used in the update to RTTvar is its value before updating sRTT
147 itself using the second assignment. That is, updating RTTvar and sRTT MUST be
148 computed in the above order.
149
150 The above SHOULD be computed using alpha=1/8 and beta=1/4 (as suggested in [JK88]).
151
152 After the computation, a host MUST update:
153
154     RT0 <- sRTT + max (G, K*RTTvar)
155
156 (where K=4, from 2.2, G is only larger than RTTvar if RTTvar = 0)
157
158 a) after the transmission of the first-mentioned segment (4:30:20)
159
160     sRTT = 4s
161     RT0 = 8s (as the segment was resent after 8 seconds)
162     RT0 came from sRTT + 4*RTTvar =>
163     RTTvar = 1s
164
165 b) after the first retransmission (4:30:28)
166
167     RT0 doubles when it triggers (5.5), no new measures for RTT are made, so sRTT and
168     RTTvar remain the same.
169
170     sRTT = 4s
171     RT0 = 16s
172     RTTvar = 1s
173
174 c) after the reception of the first acknowledgement (4:30:30)
175
176 Assuming that the ACK came from the last-sent packet, with a new RTT measure
177 R' of 2 seconds. The algorithm updates as specified above, according to (2.3).
178
179     RTTvar <- (1 - beta) * RTTvar + beta * |sRTT - R'|
180     sRTT <- (1 - alpha) * sRTT + alpha * R'
181     RT0 <- sRTT + max (G, K*RTTvar)
182
183     RTTvar <- (1 - 1/4) * 1 + 1/4 * |4 - 2| = 3/4 + 2/4 = 5/4 = 1.25s
184     sRTT <- (1 - 1/8) * 4 + 1/8 * 2 = 7/8 * 4 + 2/8 = 30/8 = 3.75s
185     RT0 <- 3.75 + 4*1.25 = 8.75s
186
187 d) after the reception of the second acknowledgement (4:30:32)
188
189 Same as above, R' = 2:
190
191     RTTvar <- 3/4 * 1.25 + 1/4 * |3.75 - 2| = 1.375s
192     sRTT <- 7/8 * 3.75 + 1/8 * 2 = 3.53125s
193     RT0 <- 3.53125 + 4*1.375 = 9.03125s
194
195 5. TCP 3
196
197 Stations A and B are connected via a 100Mbps link between the Earth and a communication
198 satellite at an altitude of 36,000 km. Assume that the signal propagation speed equals
199 the speed of light (300000km/s).
200
201 a) Calculate the minimum round-trip time (RTT) for the link.

```

202
203 The connection must travel to the satellite and back twice, that is $4 \times 36,000 \text{ km} =$
204 $= 144,000 \text{ km}$, to get back, this gives a RTT of $144,000/300,000 = 0.48 \text{ seconds}$.
205
206 b) Calculate the bandwidth-delay product of the link. Explain the meaning of this
207 product.
208
209 $\text{BPD} = 100\text{Mbps} \times 0.48 \text{ seconds} = 48\text{Mb} = 6\text{MB}$
210
211 This is the highest amount of traffic that reside inside in the network at any time.
212 That is, packets that have been sent but not yet recieved.
213
214 c) What's the minimum time to transfer a 25MB file from A to B? Include the connection
215 establishment time. Consider the transfer finished when the last ACK has arrived at
216 the sender. Assume that there are no losses. Use the RTT value computed above.
217
218 First, a three-way connection has to be established. This takes RTT time, as the
219 sender will not start transmitting before the first SYN-ACK has been recieved.
220
221 After the connection has been established, we can push things into the pipe at 100 Mbit
222 $= 12.5\text{MB per second}$. To send a 25MB file this takes about 2 seconds. After the final
223 byte has been pushed into the pipe, we'll need to wait for the final ACK, which will
224 take RTT time as well.
225
226 Total time: $2\text{s} + 2 \times \text{RTT} = 2 + 2 \times 0.48 = 2.96 \text{ seconds}$.